

IF2130 – Organisasi dan Arsitektur Komputer

sumber: Greg Kesden, CMU 15-213, 2012

Machine-Level Programming: Memory Layout dan Buffer Overflow

Achmad Imam Kistijantoro (imam@informatika.org)

Anggrahita Bayu Sasmita

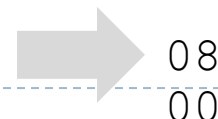
Rahmat Mulyawan

Infal Syafalni

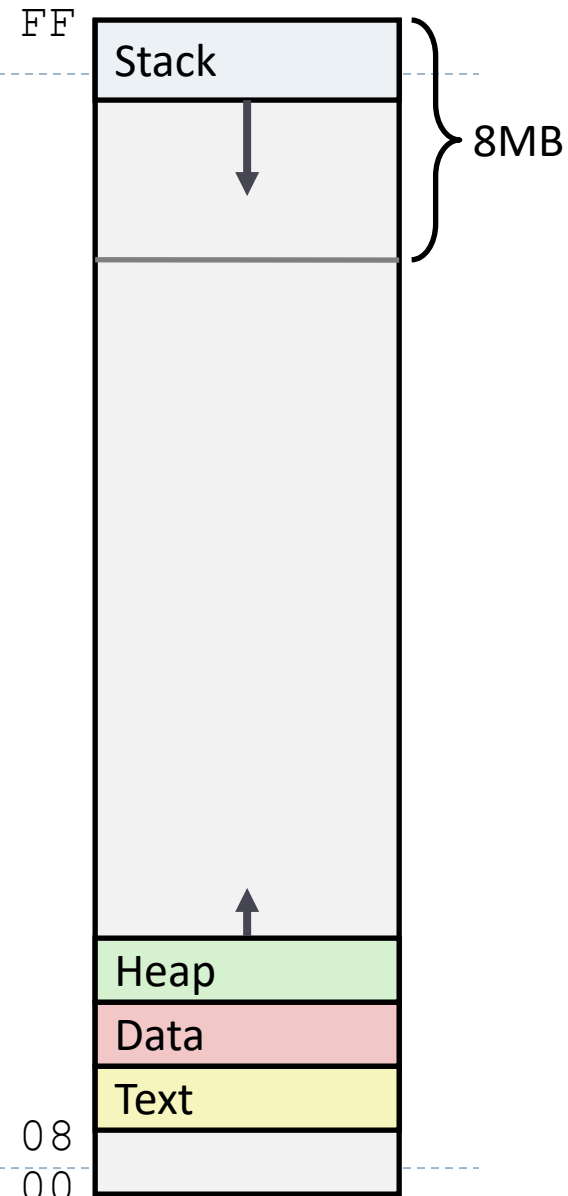
IA32 Linux Memory Layout

- ▶ **Stack**
 - ▶ Runtime stack (8MB limit)
 - ▶ E. g., local variables
- ▶ **Heap**
 - ▶ Dynamically allocated storage
 - ▶ When call `malloc()`, `calloc()`, `new()`
- ▶ **Data**
 - ▶ Statically allocated data
 - ▶ E.g., arrays & strings declared in code
- ▶ **Text**
 - ▶ Executable machine instructions
 - ▶ Read-only

Upper 2 hex digits
= 8 bits of address



not drawn to scale



Memory Allocation Example

not drawn to scale

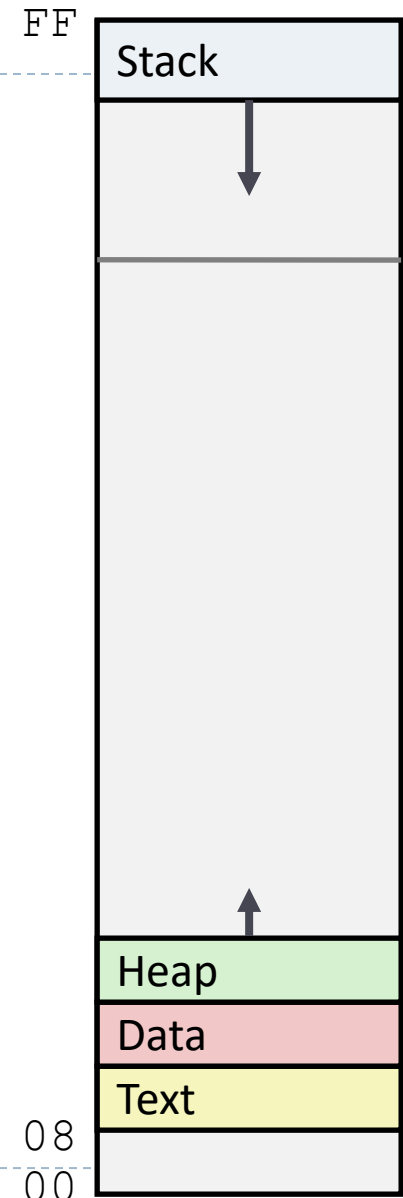
```
char big_array[1<<24]; /* 16 MB */
char huge_array[1<<28]; /* 256 MB */

int beyond;
char *p1, *p2, *p3, *p4;

int useless() { return 0; }

int main()
{
    p1 = malloc(1 << 28); /* 256 MB */
    p2 = malloc(1 << 8); /* 256 B */
    p3 = malloc(1 << 28); /* 256 MB */
    p4 = malloc(1 << 8); /* 256 B */
    /* Some print statements ... */
}
```

Where does everything go?



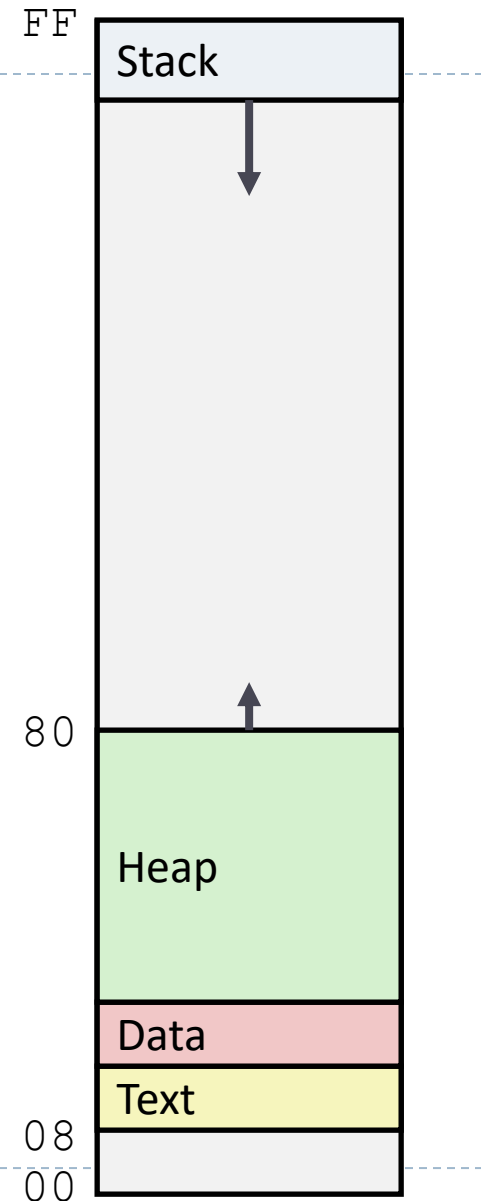
IA32 Example Addresses

address range $\sim 2^{32}$

| | |
|----------------|------------|
| \$esp | 0xffffbcd0 |
| p3 | 0x65586008 |
| p1 | 0x55585008 |
| p4 | 0x1904a110 |
| p2 | 0x1904a008 |
| &p2 | 0x18049760 |
| &beyond | 0x08049744 |
| big_array | 0x18049780 |
| huge_array | 0x08049760 |
| main() | 0x080483c6 |
| useless() | 0x08049744 |
| final malloc() | 0x006be166 |

malloc() is dynamically linked
address determined at runtime

not drawn to scale



```
imam@DELL-2020:~/if2130$ more memory_alloc.c
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
char big_array[1<<24];
```

```
char huge_array[1<<28];
```

```
int beyond;
```

```
char *p1, *p2, *p3, *p4;
```

```
int useless() { return 0; }
```

```
int main() {
```

```
    p1 = malloc(1<<28);
```

```
    p2 = malloc(1<<8);
```

```
    p3 = malloc(1<<28);
```

```
    p4 = malloc(1<<8);
```

```
    register long i asm("rsp");
```

```
    printf("rsp : %#010lx\n", i);
```

```
    printf("p3 : %#010lx\n", (long) p3);
```

```
    printf("p1 : %#010lx\n", (long) p1);
```

```
    printf("p4 : %#010lx\n", (long) p4);
```

```
    printf("p2 : %#010lx\n", (long) p2);
```

```
    printf("&p2 : %#010lx\n", (long) &p2);
```

```
    printf("&beyond: %#010lx\n", (long) &beyond);
```

```
    printf("big_array : %#010lx\n", (long) big_array);
```

```
    printf("huge_array : %#010lx\n", (long) huge_array);
```

```
    printf("main() : %#010lx\n", (long) main);
```

```
    printf("useless() : %#010lx\n", (long) useless);
```

```
    printf("malloc() : %#010lx\n", (long) malloc);
```

```
    getchar();
```

```
}
```

```
imam@DELL-2020:~/if2130$ ./memory_alloc
```

```
rsp      : 0x7ffe71da4d10
p3       : 0x7f72940cd010
p1       : 0x7f72a40ce010
p4       : 0x55da49fa03b0
p2       : 0x55da49fa02a0
&p2      : 0x55da38981040
&beyond  : 0x55da49981080
big_array : 0x55da48981080
huge_array : 0x55da38981060
main()   : 0x55da3897e188
useless() : 0x55da3897e179
malloc() : 0x7f72b416c260
```

```
imam@DELL-2020:~/if2130$ more /proc/412/maps
```

```
55da3897d000-55da3897e000 r--p 00000000 08:10 483032 /home/imam/if2130/memory_alloc
55da3897e000-55da3897f000 r-xp 00001000 08:10 483032 /home/imam/if2130/memory_alloc
55da3897f000-55da38980000 r--p 00002000 08:10 483032 /home/imam/if2130/memory_alloc
55da38980000-55da38981000 r--p 00002000 08:10 483032 /home/imam/if2130/memory_alloc
55da38981000-55da38982000 rw-p 00003000 08:10 483032 /home/imam/if2130/memory_alloc
55da38982000-55da49982000 rw-p 00000000 00:00 0
55da49fa0000-55da49fc1000 rw-p 00000000 00:00 0
7f72940cd000-7f72b40cf000 rw-p 00000000 00:00 0
7f72b40cf000-7f72b40f4000 r--p 00000000 08:10 30676 /lib/x86_64-linux-gnu/libc-2.31.so
7f72b40f4000-7f72b426c000 r-xp 00025000 08:10 30676 /lib/x86_64-linux-gnu/libc-2.31.so
7f72b426c000-7f72b42b6000 r--p 0019d000 08:10 30676 /lib/x86_64-linux-gnu/libc-2.31.so
7f72b42b6000-7f72b42b7000 ---p 001e7000 08:10 30676 /lib/x86_64-linux-gnu/libc-2.31.so
7f72b42b7000-7f72b42ba000 r--p 001e7000 08:10 30676 /lib/x86_64-linux-gnu/libc-2.31.so
7f72b42ba000-7f72b42bd000 rw-p 001ea000 08:10 30676 /lib/x86_64-linux-gnu/libc-2.31.so
7f72b42bd000-7f72b42c3000 rw-p 00000000 00:00 0
7f72b42d3000-7f72b42d4000 r--p 00000000 08:10 30668 /lib/x86_64-linux-gnu/ld-2.31.so
7f72b42d4000-7f72b42f7000 r-xp 00001000 08:10 30668 /lib/x86_64-linux-gnu/ld-2.31.so
7f72b42f7000-7f72b42ff000 r--p 00024000 08:10 30668 /lib/x86_64-linux-gnu/ld-2.31.so
7f72b4300000-7f72b4301000 r--p 0002c000 08:10 30668 /lib/x86_64-linux-gnu/ld-2.31.so
7f72b4301000-7f72b4302000 rw-p 0002d000 08:10 30668 /lib/x86_64-linux-gnu/ld-2.31.so
7f72b4302000-7f72b4303000 rw-p 00000000 00:00 0
7ffe71d86000-7ffe71da7000 rw-p 00000000 00:00 0
7ffe71dda000-7ffe71ddd000 r--p 00000000 00:00 0
7ffe71ddd000-7ffe71ddf000 r-xp 00000000 00:00 0
```

[heap]

[stack]

[vvar]

[vdso]

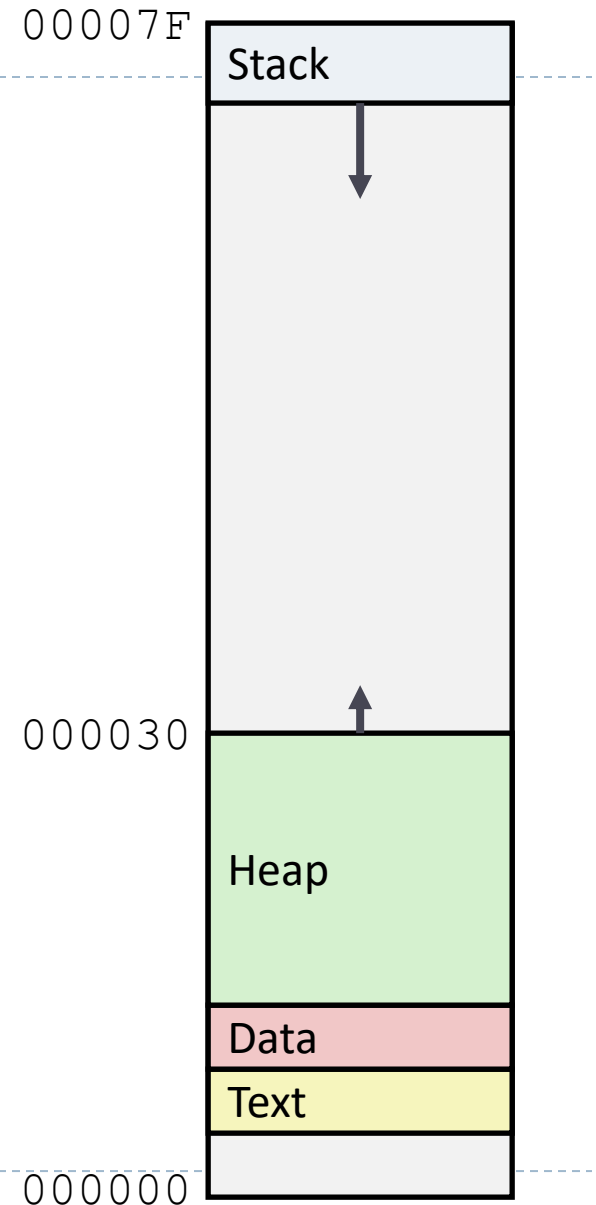
x86-64 Example Addresses

address range $\sim 2^{47}$

| | |
|-----------------------------|----------------------------------|
| <code>\$rsp</code> | <code>0x00007fffffff8d1f8</code> |
| <code>p3</code> | <code>0x00002aaabaadd010</code> |
| <code>p1</code> | <code>0x00002aaaaadc010</code> |
| <code>p4</code> | <code>0x0000000011501120</code> |
| <code>p2</code> | <code>0x0000000011501010</code> |
| <code>&p2</code> | <code>0x0000000010500a60</code> |
| <code>&beyond</code> | <code>0x0000000000500a44</code> |
| <code>big_array</code> | <code>0x0000000010500a80</code> |
| <code>huge_array</code> | <code>0x0000000000500a50</code> |
| <code>main()</code> | <code>0x0000000000400510</code> |
| <code>useless()</code> | <code>0x0000000000400500</code> |
| <code>final malloc()</code> | <code>0x000000386ae6a170</code> |

`malloc()` is dynamically linked
address determined at runtime

not drawn to scale



Today

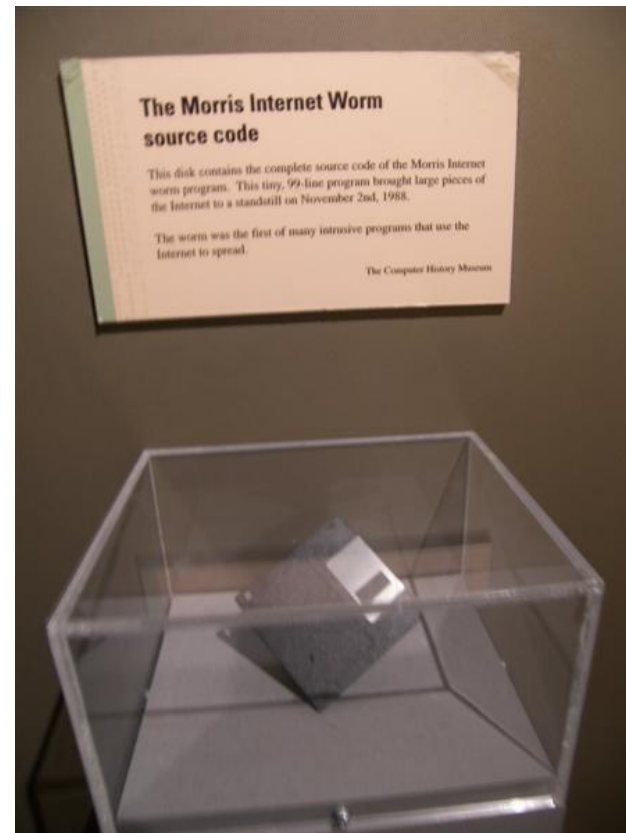
■ Structures

- Alignment
- ▶ Unions
- ▶ Memory Layout
- ▶ **Buffer Overflow**
 - ▶ Vulnerability
 - ▶ Protection



Internet Worm and IM War

- ▶ **November, 1988**
 - ▶ Internet Worm attacks thousands of Internet hosts.
 - ▶ How did it happen?



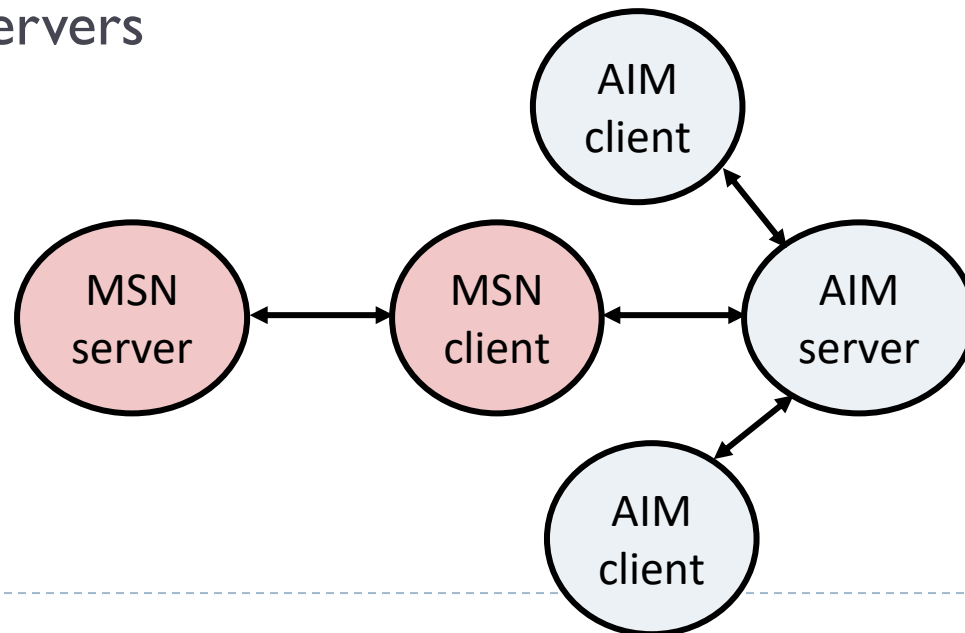
Internet Worm and IM War

▶ November, 1988

- ▶ Internet Worm attacks thousands of Internet hosts.
- ▶ How did it happen?

▶ July, 1999

- ▶ Microsoft launches MSN Messenger (instant messaging system).
- ▶ Messenger clients can access popular AOL Instant Messaging Service (AIM) servers



Internet Worm and IM War (cont.)

▶ August 1999

- ▶ Mysteriously, Messenger clients can no longer access AIM servers.
 - ▶ Microsoft and AOL begin the IM war:
 - ▶ AOL changes server to disallow Messenger clients
 - ▶ Microsoft makes changes to clients to defeat AOL changes.
 - ▶ At least 13 such skirmishes.
 - ▶ How did it happen?
 - ▶ The Internet Worm and AOL/Microsoft War were both based on *stack buffer overflow* exploits!
 - ▶ many library functions do not check argument sizes.
 - ▶ allows target buffers to overflow.
-



String Library Code

► Implementation of Unix function `gets()`

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- No way to specify limit on number of characters to read
- Similar problems with other library functions
 - **`strcpy`, `strcat`**: Copy strings of arbitrary length
 - **`scanf`, `fscanf`, `sscanf`**, when given **`%s`** conversion specification

Vulnerable Buffer Code

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
void call_echo() {  
    echo();  
}
```

```
unix>./bufdemo  
Type a string:1234567  
1234567
```

```
unix>./bufdemo  
Type a string:12345678  
Segmentation Fault
```

```
unix>./bufdemo  
Type a string:123456789ABC  
Segmentation Fault
```



Buffer Overflow Disassembly

echo:

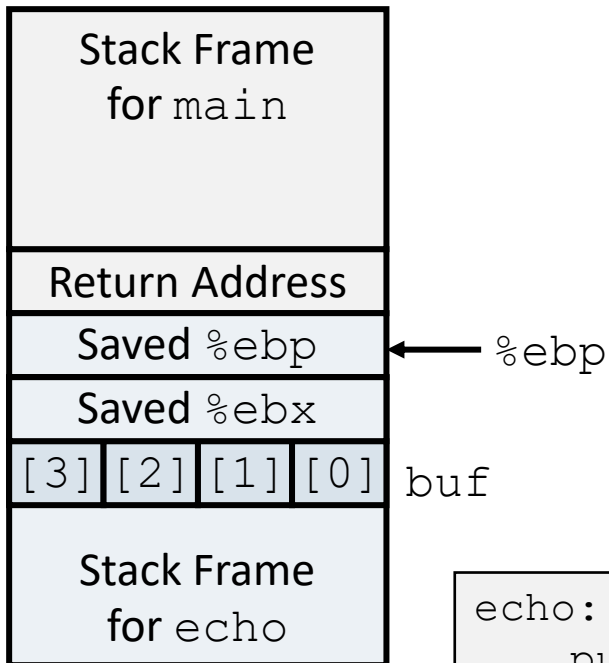
| | | | |
|----------|----------------|------|------------------------|
| 80485c5: | 55 | push | %ebp |
| 80485c6: | 89 e5 | mov | %esp,%ebp |
| 80485c8: | 53 | push | %ebx |
| 80485c9: | 83 ec 14 | sub | \$0x14,%esp |
| 80485cc: | 8d 5d f8 | lea | 0xffffffff8(%ebp),%ebx |
| 80485cf: | 89 1c 24 | mov | %ebx, (%esp) |
| 80485d2: | e8 9e ff ff ff | call | 8048575 <gets> |
| 80485d7: | 89 1c 24 | mov | %ebx, (%esp) |
| 80485da: | e8 05 fe ff ff | call | 80483e4 <puts@plt> |
| 80485df: | 83 c4 14 | add | \$0x14,%esp |
| 80485e2: | 5b | pop | %ebx |
| 80485e3: | 5d | pop | %ebp |
| 80485e4: | c3 | ret | |

call_echo:

| | | | |
|----------|----------------|-------|----------------|
| 80485eb: | e8 d5 ff ff ff | call | 80485c5 <echo> |
| 80485f0: | c9 | leave | |
| 80485f1: | c3 | ret | |

Buffer Overflow Stack

Before call to gets



```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

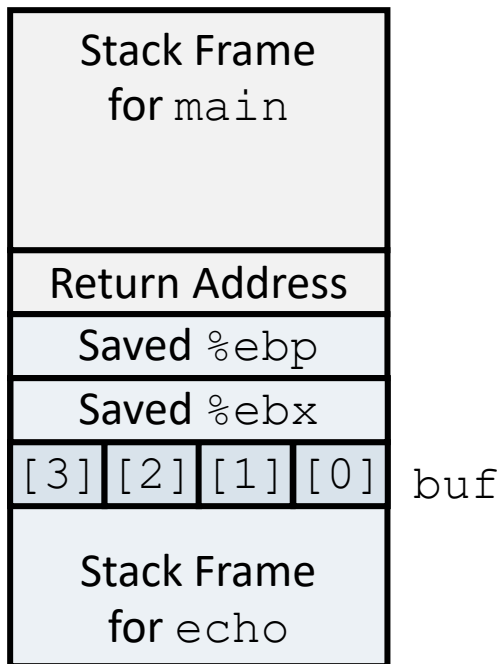
```
echo:
    pushl %ebp                # Save %ebp on stack
    movl %esp, %ebp
    pushl %ebx                # Save %ebx
    subl $20, %esp            # Allocate stack space
    leal -8(%ebp), %ebx        # Compute buf as %ebp-8
    movl %ebx, (%esp)          # Push buf on stack
    call gets                  # Call gets
    . . .
```

Buffer Overflow Stack Example

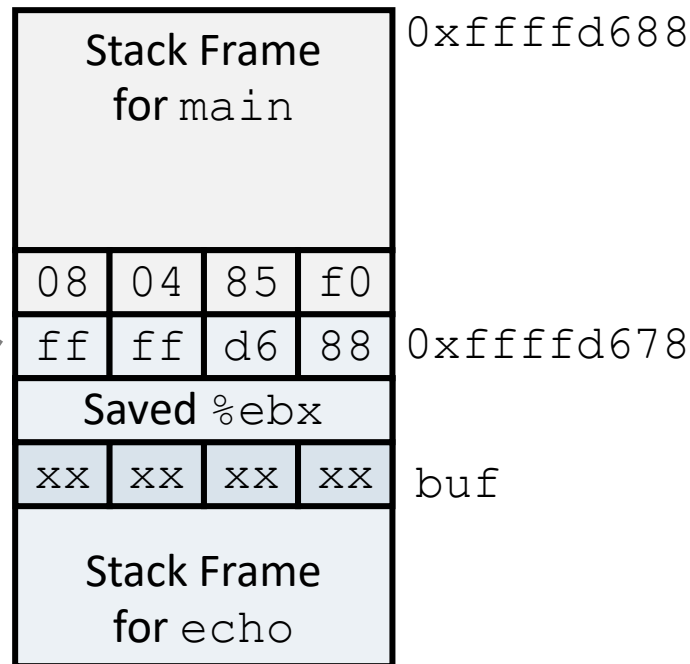
```

unix> gdb bufdemo
(gdb) break echo
Breakpoint 1 at 0x80485c9
(gdb) run
Breakpoint 1, 0x80485c9 in echo ()
(gdb) print /x $ebp
$1 = 0xffffd678
(gdb) print /x *(unsigned *)$ebp
$2 = 0xffffd688
(gdb) print /x *((unsigned *)$ebp + 1)
$3 = 0x80485f0
  
```

Before call to gets



Before call to gets

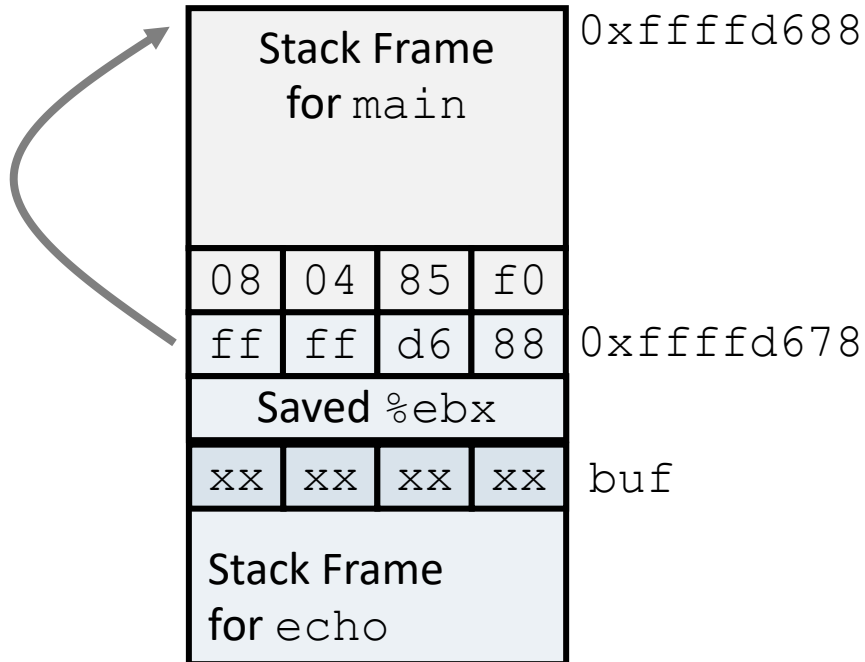


```

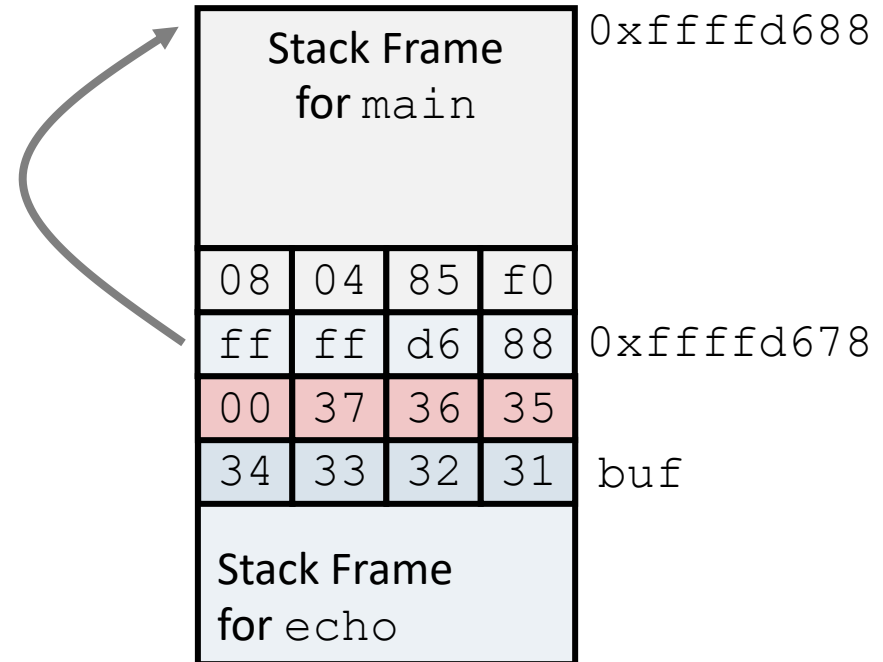
80485eb:  e8 d5 ff ff ff  call  80485c5 <echo>
80485f0:  c9                leave
  
```


Buffer Overflow Example #1

Before call to gets



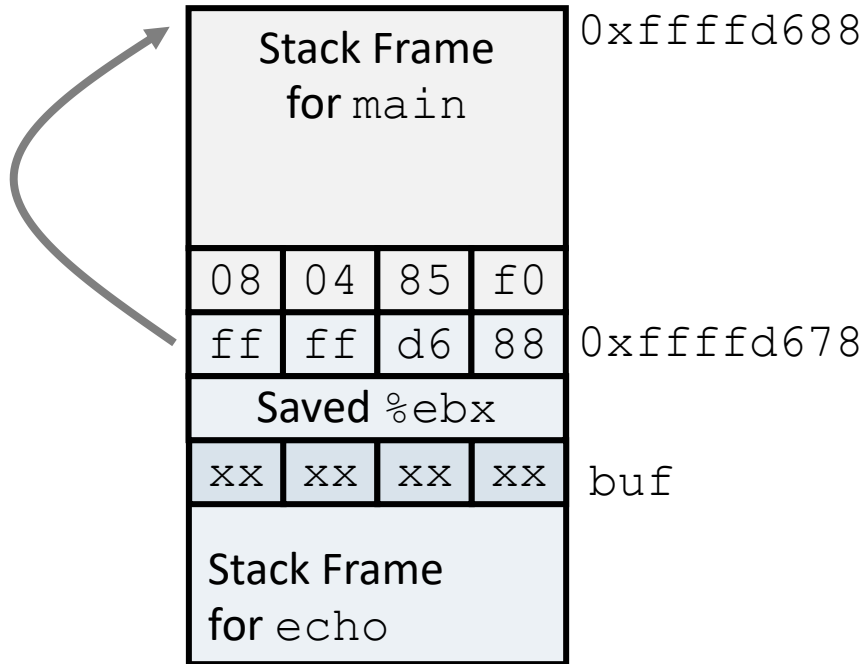
Input 1234567



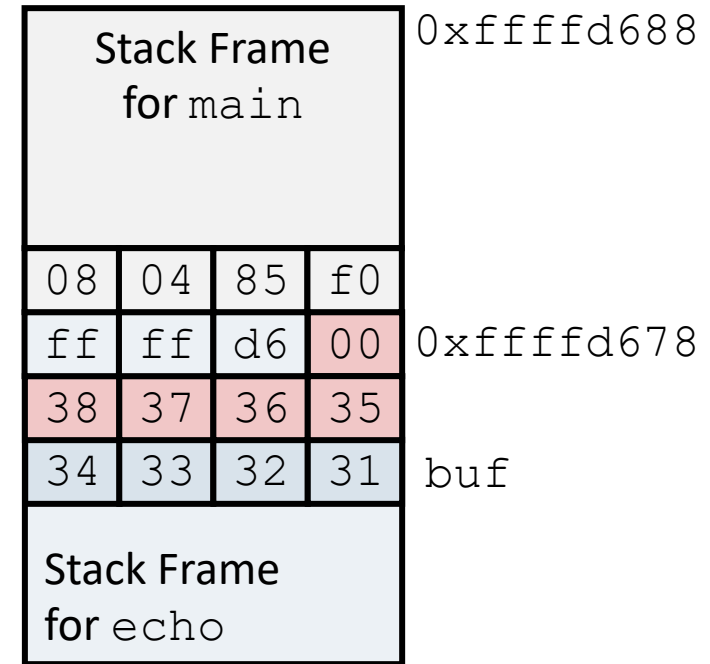
Overflow buf, and corrupt %ebx,
but no problem

Buffer Overflow Example #2

Before call to gets



Input 12345678

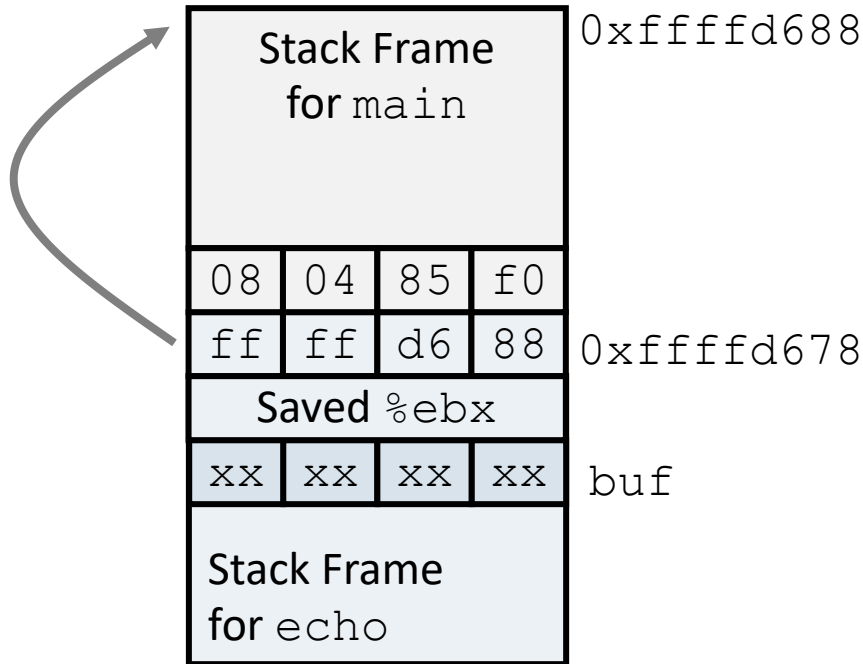


Base pointer corrupted

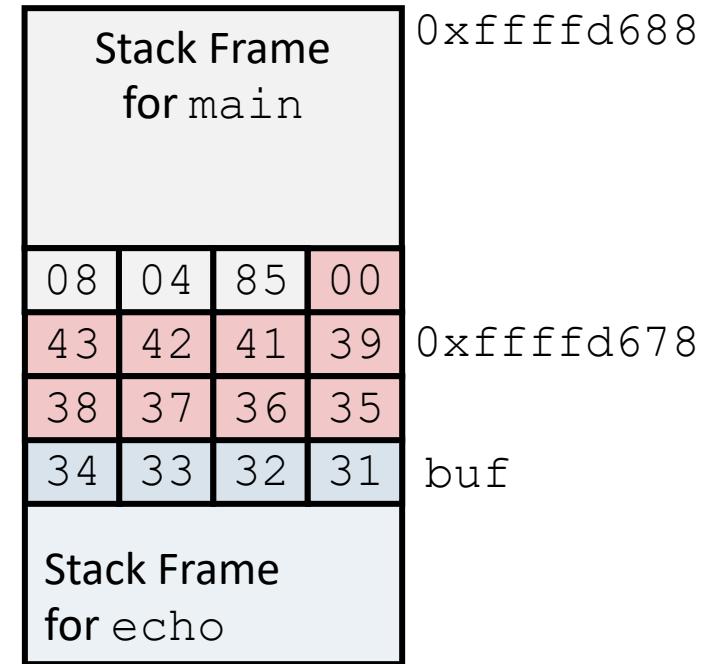
```
. . .
80485eb:  e8 d5 ff ff ff  call    80485c5 <echo>
80485f0:  c9              leave   # Set %ebp to corrupted value
80485f1:  c3             ret
```

Buffer Overflow Example #3

Before call to gets



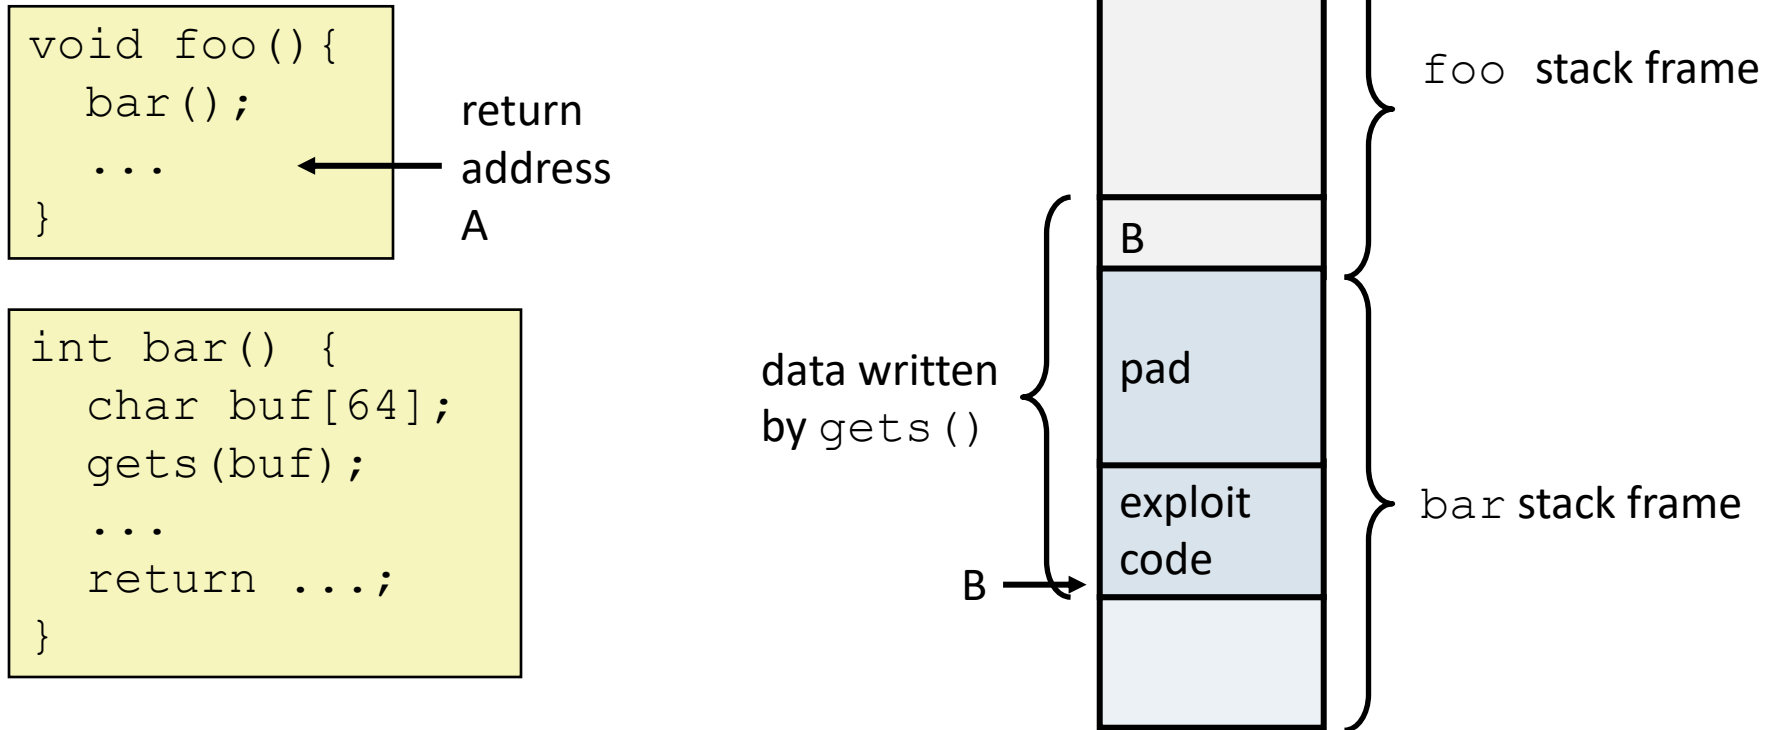
Input 123456789ABC



Return address corrupted

```
80485eb:  e8 d5 ff ff ff  call  80485c5 <echo>
80485f0:  c9              leave # Desired return point
```

Malicious Use of Buffer Overflow



- ▶ Input string contains byte representation of executable code
- ▶ Overwrite return address A with address of buffer B
- ▶ When bar() executes ret, will jump to exploit code

Exploits Based on Buffer Overflows

- ▶ *Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines*
- ▶ Internet worm
 - ▶ Early versions of the finger server (fingerd) used `gets()` to read the argument sent by the client:
 - ▶ `finger droh@cs.cmu.edu`
 - ▶ Worm attacked fingerd server by sending phony argument:
 - ▶ `finger "exploit-code padding new-return-address"`
 - ▶ exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.



Exploits Based on Buffer Overflows

- ▶ *Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines*
- ▶ IMWar
 - ▶ AOL exploited existing buffer overflow bug in AIM clients
 - ▶ exploit code: returned 4-byte signature (the bytes at some location in the AIM client) to server.
 - ▶ When Microsoft changed code to match signature, AOL changed signature location.



Date: Wed, 11 Aug 1999 11:30:57 -0700 (PDT)
From: Phil Bucking <philbucking@yahoo.com>
Subject: AOL exploiting buffer overrun bug in their own software!
To: rms@pharlap.com

Mr. Smith,

I am writing you because I have discovered something that I think you might find interesting because you are an Internet security expert with experience in this area. I have also tried to contact AOL but received no response.

I am a developer who has been working on a revolutionary new instant messaging client that should be released later this year.

...
It appears that the AIM client has a buffer overrun bug. By itself this might not be the end of the world, as MS surely has had its share. But AOL is now *exploiting their own buffer overrun bug* to help in its efforts to block MS Instant Messenger.

....
Since you have significant credibility with the press I hope that you can use this information to help inform people that behind AOL's friendly exterior they are nefariously compromising peoples' security.

Sincerely,
Phil Bucking
Founder, Bucking Consulting
philbucking@yahoo.com

*It was later determined that this email
originated from within Microsoft!*



Avoiding Overflow Vulnerability

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    fgets(buf, 4, stdin);  
    puts(buf);  
}
```

- ▶ Use library routines that limit string lengths
 - ▶ **fgets** instead of **gets**
 - ▶ **strncpy** instead of **strcpy**
 - ▶ Don't use **scanf** with **%s** conversion specification
 - ▶ Use **fgets** to read the string
 - ▶ Or use **%ns** where **n** is a suitable integer



System-Level Protections

▶ Randomized stack offsets

- ▶ At start of program, allocate random amount of space on stack
- ▶ Makes it difficult for hacker to predict beginning of inserted code

▶ Nonexecutable code segments

- ▶ In traditional x86, can mark region of memory as either “read-only” or “writeable”
 - ▶ Can execute anything readable
- ▶ X86-64 added explicit “execute” permission

```
unix> gdb bufdemo
(gdb) break echo

(gdb) run
(gdb) print /x $ebp
$1 = 0xfffffc638

(gdb) run
(gdb) print /x $ebp
$2 = 0xffffbb08

(gdb) run
(gdb) print /x $ebp
$3 = 0xffffc6a8
```

Stack Canaries

▶ Idea

- ▶ Place special value (“canary”) on stack just beyond buffer
- ▶ Check for corruption before exiting function

▶ GCC Implementation

- ▶ **-fstack-protector**
- ▶ **-fstack-protector-all**

```
unix>./bufdemo-protected  
Type a string:1234  
1234
```

```
unix>./bufdemo-protected  
Type a string:12345  
*** stack smashing detected ***
```



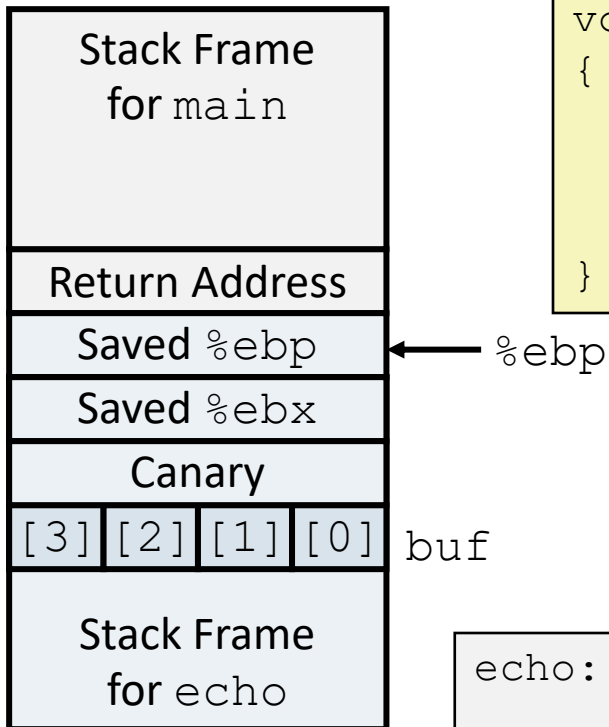
Protected Buffer Disassembly

echo:

| | | | |
|----------|----------------------|------|-------------------------|
| 804864d: | 55 | push | %ebp |
| 804864e: | 89 e5 | mov | %esp,%ebp |
| 8048650: | 53 | push | %ebx |
| 8048651: | 83 ec 14 | sub | \$0x14,%esp |
| 8048654: | 65 a1 14 00 00 00 | mov | %gs:0x14,%eax |
| 804865a: | 89 45 f8 | mov | %eax,0xffffffff8(%ebp) |
| 804865d: | 31 c0 | xor | %eax,%eax |
| 804865f: | 8d 5d f4 | lea | 0xfffffffff4(%ebp),%ebx |
| 8048662: | 89 1c 24 | mov | %ebx, (%esp) |
| 8048665: | e8 77 ff ff ff | call | 80485e1 <gets> |
| 804866a: | 89 1c 24 | mov | %ebx, (%esp) |
| 804866d: | e8 ca fd ff ff | call | 804843c <puts@plt> |
| 8048672: | 8b 45 f8 | mov | 0xfffffffff8(%ebp),%eax |
| 8048675: | 65 33 05 14 00 00 00 | xor | %gs:0x14,%eax |
| 804867c: | 74 05 | je | 8048683 <echo+0x36> |
| 804867e: | e8 a9 fd ff ff | call | 804842c <FAIL> |
| 8048683: | 83 c4 14 | add | \$0x14,%esp |
| 8048686: | 5b | pop | %ebx |
| 8048687: | 5d | pop | %ebp |
| 8048688: | c3 | ret | |

Setting Up Canary

Before call to gets

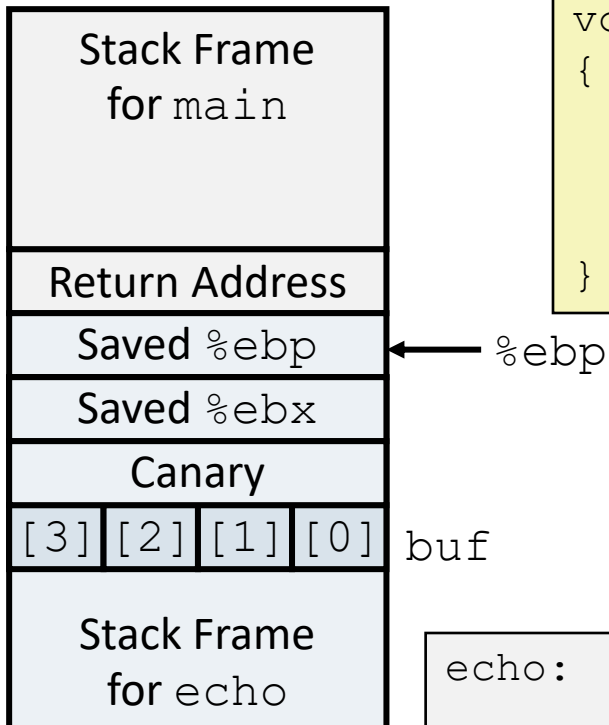


```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    . . .  
    movl    %gs:20, %eax    # Get canary  
    movl    %eax, -8(%ebp)  # Put on stack  
    xorl    %eax, %eax     # Erase canary  
    . . .
```

Checking Canary

Before call to gets

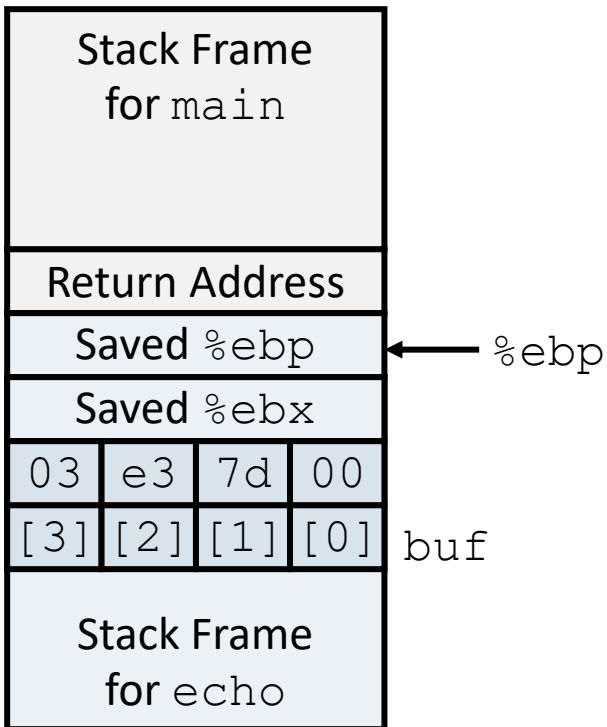


```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

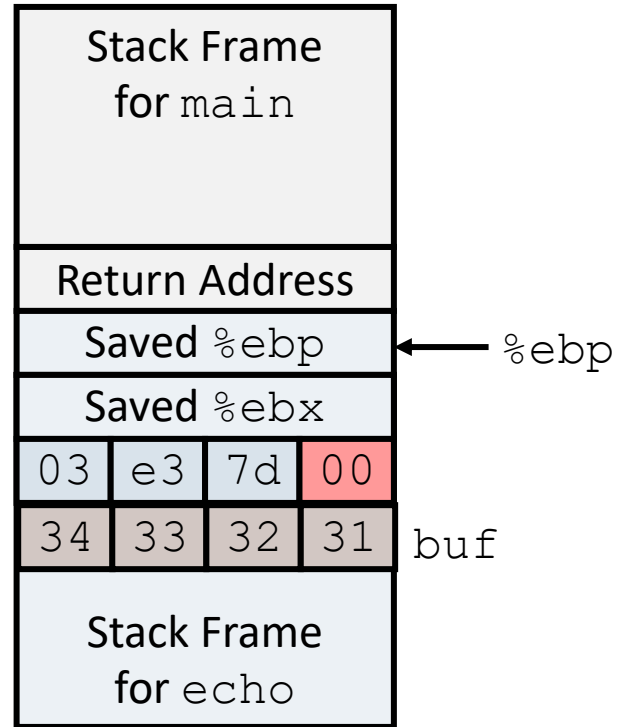
```
echo:
    . . .
    movl    -8(%ebp), %eax    # Retrieve from stack
    xorl    %gs:20, %eax      # Compare with Canary
    je      .L24              # Same: skip ahead
    call    __stack_chk_fail  # ERROR
.L24:
    . . .
```

Canary Example

Before call to gets



Input 1234



```
(gdb) break echo
(gdb) run
(gdb) stepi 3
(gdb) print /x *((unsigned *) $ebp - 2)
$1 = 0x3e37d00
```

Benign corruption!
(allows programmers to make
silent off-by-one errors)

Worms and Viruses

- ▶ **Worm:** A program that
 - ▶ Can run by itself
 - ▶ Can propagate a fully working version of itself to other computers
- ▶ **Virus:** Code that
 - ▶ Add itself to other programs
 - ▶ Cannot run independently
- ▶ Both are (usually) designed to spread among computers and to wreak havoc



Today

■ Structures

- Alignment

▶ Unions

▶ Memory Layout

▶ Buffer Overflow

- ▶ Vulnerability
- ▶ Protection

